

HLS4ML Layer Implementation Tutorial

NOTE: THIS TUTORIAL WILL ONLY TEACH YOU HOW TO IMPLEMENT TF.KERAS LAYERS & CUSTOM LAYER (USING PR 528) IN HLS4ML. FOR PYTORCH LAYER & OTHER LAYERS PLEASE DISCUSS WITH YOUR ADVISOR

Regular TensorFlow Keras Layer Implementation Overview

1. Please check the current supported Layers at <https://github.com/fastmachinelearning/hls4ml/tree/main/hls4ml/converters>
2. Transformers and Bidirectional layers are also available based on the request.
3. The implementation for regular TensorFlow Keras Layer contains both the Python Files and the C++ file.

Python Files

1. Adding a converter under this directory
 - a. <https://github.com/fastmachinelearning/hls4ml/tree/main/hls4ml/converters>
 - b. The converter parses the needed args from the Keras layer to a new layer object.
 - c. The converter also defines the output shape of the layer.
 - d. After the converter is added, you should see the correct output shape printed by pushing the model through HLS4ML.
2. Adding required implementation in layers.py
 - a. <https://github.com/fastmachinelearning/hls4ml/blob/main/hls4ml/model/layers.py>
 - b. In layers.py, you would have to define the attribute of the current layer.
 - c. You'll also need to read the weight & bias from the model. The path & name of the weight & bias has to be exactly the same as in the saved model. You can use HDFview <https://www.hdfgroup.org/downloads/hdfview/> to see the weight/bias path & name in the saved model.
 - d. Last, you need to add the layer into the layer map <https://github.com/fastmachinelearning/hls4ml/blob/main/hls4ml/model/layers.py#L1256>
3. Adding a python template under this directory
 - a. <https://github.com/fastmachinelearning/hls4ml/tree/main/hls4ml/backends/vivado/passes>
 - b. In this file, you'll need to first define the config_template for all the needed layers. For example, if your layer contains both A and B layers, and the B layer contains C & D layers. you'll need to define the config_template for A,B,C & D layer. This config_template will generate the config for each layer, and these configs will be passed into the C++ functions (will be defined in the C++ file).
 - c. After this file is implemented, you should be able to see all layer configs in the *parameter.h* file.

C++ File

1. Adding HLS functions under this directory

- https://github.com/fastmachinelearning/hls4ml/tree/main/hls4ml/templates/vivado/nnet_utils
- In this file, you'll need to transfer all needed functions from Keras API (python) to C++ to reach the functional correct behavior.
- You don't have to worry about the optimization at the beginning. Make sure the layer is working properly and then you can think about the optimization.

Example of adding a normalization layer in HLS4ML

First, find the normalization layer in Keras:

[LayerNormalization layer \(keras.io\)](https://keras.io/layers/normalization/)

- Read the description and make sure you fully understand the structure of the layer.
- Test your understanding by building a custom layer in Keras. Give the same input to the Keras layer, and your custom layer. Make sure the output is the same.

For example:

- Create a data set that needs to be fed into the model:

```
#2batch*(3,4)
x=x.reshape((2,3,4))
x

array([[[-0.9          , -0.82173913, -0.74347826, -0.66521739],
        [-0.58695652, -0.50869565, -0.43043478, -0.35217391],
        [-0.27391304, -0.19565217, -0.1173913 , -0.03913043]],
       [[ 0.03913043,  0.1173913 ,  0.19565217,  0.27391304],
        [ 0.35217391,  0.43043478,  0.50869565,  0.58695652],
        [ 0.66521739,  0.74347826,  0.82173913,  0.9          ]]])
```

Here, we have 2 batches, each batch contains a set of data with dimensions 3,4.

- Create a Keras Normalization layer:

```
inputs = keras.Input(shape=x.shape[1:])

output = tf.keras.layers.LayerNormalization(epsilon=1e-6)(inputs)
model = keras.Model(inputs, output)
```

Each time we will feed one batch to the model, so the input shape is (3,4). “3” is the sequence length, and “4” is the feature dimension.

By default, LayerNormalization layer will normalize the data on the last axis.

- Then, we can use keras prediction to find the output of the Keras model:

```

y_keras = model.predict(x)
print("y_keras prediction is:")
print(y_keras)

1/1 [=====] - 0s 58ms/step
y_keras prediction is:
[[[-1.3415527  -0.44718456  0.44718456  1.3415532 ]
  [-1.3415532  -0.44718456  0.4471841  1.3415532 ]
  [-1.3415532  -0.44718432  0.44718444  1.3415532 ]

  [-1.3415532  -0.44718444  0.44718432  1.3415532 ]
  [-1.3415532  -0.4471841  0.44718456  1.3415532 ]
  [-1.3415542  -0.44718552  0.4471836  1.3415518 ]]]

```

4) Implement a function/custom layer to verify your understanding of the model:

```

k=4 #k is the feature dim
x_i_normalized = np.zeros(4)
mean_i = sum(test[j] for j in range(k)) / k
var_i = sum((test[j] - mean_i) ** 2 for j in range(k)) / k
for j in range(k):
    x_i_normalized[j] = (test[j] - mean_i) / np.sqrt(var_i)
    print("data_dff is :", (test[j] - mean_i))
print("table_inv is :", 1/np.sqrt(var_i))
x_i_normalized|

array([-1.34164079, -0.4472136 ,  0.4472136 ,  1.34164079])

```

Here we verified that our understanding of the model matches the Keras output.
Thus, we can start to implement the model in hls4ml.

3. Implement the Converter (Parser) in hls4ml

In [hls4ml/hls4ml/converter/keras at main · fastmachinelearning/hls4ml · GitHub](https://github.com/fastmachinelearning/hls4ml/blob/main/hls4ml/converter/keras.py), we can find a lot of examples of the parser for each layer. As you can see, the task of the parser is to extract the configuration parameters from the layer, such as the number of inputs, the number of feature dimensions, the sequence length, the attention axis, etc...

The parser of each layer has a similar structure, with the inputs, such as `keras_layer`, `input_names`, `input_shapes`, `data_reader`, `config`, and outputs, such as `layer`, `[shape for shape in input_shapes[0]]`.

[The parser of the normalization layer:](#)

```

135 @keras_handler('LayerNormalization')
136 def parse_layernorm_layer(keras_layer, input_names, input_shapes, data_reader, config):
137     assert('LayerNormalization' in keras_layer['class_name'])
138
139     layer = parse_default_keras_layer(keras_layer, input_names)
140
141     in_size = 1
142     for dim in input_shapes[0][1:]:
143         in_size *= dim
144
145     layer['axis'] = keras_layer['config']['axis'] if (keras_layer['config']['axis'][0]==2) else False
146     if layer['axis'] is False:
147         raise Exception('assigning the axis is not currently supported by hls4ml, only axis 2 is supported')
148
149     if not((len(input_shapes[0])) == 3 ):
150         raise Exception('input size is not currently supported by hls4ml, only dim3 is supported')
151     if len(input_shapes[0])==3:
152         layer['seq_len'] = input_shapes[0][-2]
153     else: layer['seq_len'] = 1
154     layer['n_in'] = in_size
155     layer['n_out'] = layer['n_in']
156
157     return layer, [shape for shape in input_shapes[0]]

```

When implementing the parser, we need to think about what parameters are needed for the layer. First, we need the number of the input for the layer. For the 1-d model, the input size would be the dimension of the last axis. However, when the model is 2-d or 3d, the input dimension would be the product of each axis' dimension (except batch).

Second, We need to know which axis will be used to calculate the normalized value. If the axis is not to be considered supported in hls4ml, we need to raise an assertion. For the current hls4ml normalization layer, we only supported 2-d inputs with normalization on the second axis, so I made two assertions. Third, we need to find the number of output and sequence length.

4. In step 3, we load the configuration parameters into the hls4ml, but how about the weight parameters? In `hls4ml/model/layers.py`, we will load the weight/bias parameter into hls4ml, and store the previous configuration parameters into attributes.

[Normalization layer:](#)

```

class LayerNormalization(Layer):
    _expected_attributes = [
        Attribute('n_in'),
        # Attribute('axis', default=-1),
        Attribute('seq_len'),
        WeightAttribute('scale'),
        WeightAttribute('bias'),

        TypeAttribute('scale'),
        TypeAttribute('bias'),
    ]

    def initialize(self):
        inp = self.get_input_variable()
        shape = inp.shape
        dims = inp.dim_names
        self.add_output_variable(shape, dims)

        gamma = self.model.get_weights_data(self.name, 'gamma')
        beta = self.model.get_weights_data(self.name, 'beta')

        scale = gamma
        bias = beta

        self.add_weights_variable(name='scale', var_name='s{index}', data=scale)
        self.add_weights_variable(name='bias', var_name='b{index}', data=bias)

```

In the class `LayerNormalization(Layer)`, we create expected attributes for “number of input”, and “sequence length”. For the weights, we need to create “WeightAttributes”, the naming of the weight attribute needs to be well described for the data.

We then define a class function “`initialize(self)`” to initialize the weight parameters. To load the weight, we can use the function “`get_weights_data()`”, or “`add_weights_variable()`”. After we load the weight, we need to store the weight using the function “`add_weights_variable()`”. We also need to use “`self.add_output_variable(shape, dims)`” to add the output shape and name of each dimension.

The last part of this procedure is to add this new layer to the “`layer_map`” [dictionary](#).

5. Add the new [supported layer](#) in `hls4ml/utils/config.py`. However, the most updated version of hls4ml may change this supported layer to other places.

6. Create a template configuration in [hls4ml/hls4ml/backends/vivado/passes](#)

In the previous steps, we have defined the configuration parameters, and now we will create a file that can help us to convert the configuration into c++ code.

```

102 layernorm_config_template = """struct config{index} : nnet::layernorm_config {{
103     static const unsigned n_in = {n_in};
104     static const unsigned seq_len = {seq_len};
105     static const unsigned table_size = {table_size};
106     static const unsigned io_type = nnet::{iotype};
107     static const unsigned reuse_factor = {reuse};
108     static const bool store_weights_in_bram = false;
109     typedef {bias_t.name} bias_t;
110     typedef {scale_t.name} scale_t;
111     typedef {table_t.name} table_t;
112     template<class x_T, class y_T>
113     using product = nnet::product::{product_type}<x_T, y_T>;
114 }};\n"""
115
116 layernorm_function_template = 'nnet::layernormalize<{input_t}, {output_t}, {config}>({input}, {output}, {scale}, {bias});'
117
118 layernorm_include_list = ['nnet_utils/nnet_layernorm.h']

120 class LayerNormalizationConfigTemplate(LayerConfigTemplate):
121     def __init__(self):
122         super().__init__(LayerNormalization)
123         self.template = layernorm_config_template
124
125     def format(self, node):
126         params = self._default_config_params(node)
127         params['n_in'] = node.get_input_variable().size_cpp()
128         params['seq_len'] = node.get_attr('seq_len')
129         params['product_type'] = get_backend('vivado').product_type(node.get_input_variable().type.precision, node.get_weights('scale').type.precision)
130
131         return self.template.format(**params)
132
133 class LayerNormalizationFunctionTemplate(FunctionCallTemplate):
134     def __init__(self):
135         super().__init__(LayerNormalization, include_header=layernorm_include_list)
136         self.template = layernorm_function_template
137
138     def format(self, node):
139         params = self._default_function_params(node)
140         params['scale'] = node.get_weights('scale').name
141         params['bias'] = node.get_weights('bias').name
142
143         return self.template.format(**params)

```

https://github.com/Ethan0Jiang/hls4ml/blob/66003e7d1b49888dfbacc4d3399a8100012477d/hls4ml/backends/vivado/passes/core_templates.py#L120

From line 102 to 114, we create the config of the normalization layer that will be written into a config.h file. In this config template, the element in the {curly brackets} will be replaced by the configuration parameter, using the class `LayerNormalizationConfigTemplate()`. The `LayerNormalizationConfigTemplate` class has a `format` function to replace the argument in the template, which is a function that you need to implement. We also need to create a “`LayerNormalizationFunctionTemplate()`” class to define how we can call the normalization layer in the top .cpp file.

*Don't forget to add your new layer into the import line at the [top of the file](#).

7. (Optional) `hls4ml/backends/vivado/vivado_backend.py`

Sometimes, you might need to implement an [optimizer](#) for your layer.

In the optimizer, you can try to [reshape, or transpose](#) your weight, or you can also set a [default value](#) for the parameters.

```

@layer_optimizer(LayerNormalization)
def init_layernormalization(self, layer):
    if 'table_t' not in layer.attributes:
        layer.set_attr('table_t', NamedType(name=layer.name + '_table_t', precision=FixedPrecisionType(width=32, integer=5)))
    if 'table_size' not in layer.attributes:
        layer.set_attr('table_size', 2048)

```

Here, I create an attribute of table_size and table_t with some default value for the lookup table of nonlinear functions.

*don't forget to add your layer into import at the top of the file.

8. (Optional) resource_strategy

If you are implementing the resource strategy, you have to transpose the weight shape of the dense layer in [hls4ml/backends/vivado/passes/resource_strategy.py](https://github.com/Xilinx/hls4ml/blob/master/backends/vivado/passes/resource_strategy.py)

```

def transform(self, model, node):
    if isinstance(node, Dense):
        node.weights['weight'].data = np.transpose(node.weights['weight'].data)
    elif isinstance(node, Conv1D):
        node.weights['weight'].data = np.transpose(node.weights['weight'].data, axes=[2, 0, 1]) #(W,C,F) => (F,W,C)
    elif isinstance(node, SeparableConv1D):
        node.weights['depthwise'].data = np.transpose(node.weights['depthwise'].data, axes=[2, 0, 1]) #(W,C,F) => (F,W,C)
        node.weights['pointwise'].data = np.transpose(node.weights['pointwise'].data, axes=[2, 0, 1]) #(W,C,F) => (F,W,C)
    elif isinstance(node, Conv2D):
        node.weights['weight'].data = np.transpose(node.weights['weight'].data, axes=[3, 0, 1, 2]) #(H,W,C,F) => (F,H,W,C)
    elif isinstance(node, SeparableConv2D):
        node.weights['depthwise'].data = np.transpose(node.weights['depthwise'].data, axes=[3, 0, 1, 2]) #(H,W,C,F) => (F,H,W,C)
        node.weights['pointwise'].data = np.transpose(node.weights['pointwise'].data, axes=[3, 0, 1, 2]) #(H,W,C,F) => (F,H,W,C)
    elif isinstance(node, (LSTM, GRU)):
        node.weights['weight'].data = np.transpose(node.weights['weight'].data)
        node.weights['recurrent_weight'].data = np.transpose(node.weights['recurrent_weight'].data)
    elif isinstance(node, MultiHeadAttention):
        node.weights['key_weight'].data = np.transpose(node.weights['key_weight'].data, axes=[0, 2, 1])
        node.weights['query_weight'].data = np.transpose(node.weights['query_weight'].data, axes=[0, 2, 1])
        node.weights['value_weight'].data = np.transpose(node.weights['value_weight'].data, axes=[0, 2, 1])
        node.weights['attention_output_weight'].data = np.transpose(node.weights['attention_output_weight'].data, axes=[2, 0, 1])
    else:
        raise Exception('Unexpected layer {} with resource strategy'.format(node.class_name))

    node.set_attr('_weights_transposed', True)

```

*don't forget to add your layer into import at the top of the file.

9. Header file of your layer.

- 1) Take a look at the HLS code of different layers in "hls4ml/templates/vivado/nnet_utils/", such as dense layer, activation layer, mutiheadAttention layer, etc...
- 2) Do some research about C++ templates online.

3) Structure of a layer:

```
133 template<class data_T, class res_T, typename CONFIG_T>
134 void layernormalize(
135     data_T    data[CONFIG_T::n_in],
136     res_T     res[CONFIG_T::n_in],
137     typename CONFIG_T::scale_t scale[CONFIG_T::n_scale_bias],
138     typename CONFIG_T::bias_t  bias[CONFIG_T::n_scale_bias]
139 )
140 {
141     data_T cache;
142     static const unsigned dim = CONFIG_T::n_in/CONFIG_T::seq_len;
143
144     // Use a function_instantiate in case it helps to explicitly optimize unchanging weights/biases
145     #pragma HLS function_instantiate variable=scale,bias
146
147     // For parallel inputs:
148     // - completely partition arrays -- target fabric
149     // - if we have an unroll factor, limit number of multipliers
150     #pragma HLS PIPELINE II=CONFIG_T::reuse_factor
151
152     // #pragma HLS ARRAY_PARTITION variable=weights complete // remove this line for now, it breaks compression sometimes
153     #pragma HLS ARRAY_PARTITION variable=scale complete
154     #pragma HLS ARRAY_PARTITION variable=bias complete
155
156     for (int j=0; j < CONFIG_T::seq_len; ++j){
157         layernorm_1d<data_T, res_T, CONFIG_T>(data+(dim*j), res+(dim*j), scale, bias);
158     }
159
160
161 }
```

We will define a c++ template for the layer. The template has three type parameter,

<class data_T, class res_T, typename CONFIG_T>

The data_T is the input data type, by default it would be ap_fix<16,6>, the res_T is the output dataset, and the typename CONFIG_T is a struct that we defined in step6.

Usually there is one input and one output for the layer, and two weights. The input/output variables have been defined at step6, the FunctionTemplet.

First, we do not have to implement the detailed structure of the layer, but make sure the code from step1 to step8 can pass through the hls4ml workflow. To test that, write some very simple operation in the HLS code, for example, result = input +1. Then run the hls.compile() and hls.build() to make sure that there are no bugs and errors. If bugs occur, check the report in the prompt and try to fix it.

Once we verify the code can pass through the flow with no error, we can start to implement the layer in HLS code.

A few good resource and hint of writing the HLS code:

- 1) [HLS Pragmas \(xilinx.com\)](https://www.xilinx.com) Understand Pragma Reshape, ARRAY_PARTITION , unroll, pipeline, dataFlow, instantiate, those are the most common Pragma to use.
- 2) Check Hans' tutorial of how to write good HLS code. [Vivado HLS – Hans Giesen \(upenn.edu\)](https://www.upenn.edu)
- 3) Understand two different i/o strategy in HLS, parallel vs. hls_stream

Once finish implementing the layer, we can start to debug the layer

Custom Layer Using Extension API

In pull request 528 (<https://github.com/fastmachinelearning/hls4ml/pull/528>), HLS4ML adds an extension API that allows users to add custom layers into the package. The example can be found in `hls4ml/test/pytest/test_extensions.py`

https://github.com/fastmachinelearning/hls4ml/blob/main/test/pytest/test_extensions.py

If you want to add special layers such gaussian sampling layer, you might need to encapsulate it as a Keras layer first, an example can be found here

https://keras.io/guides/making_new_layers_and_models_via_subclassing/